

Modelo de Compilação e Execução

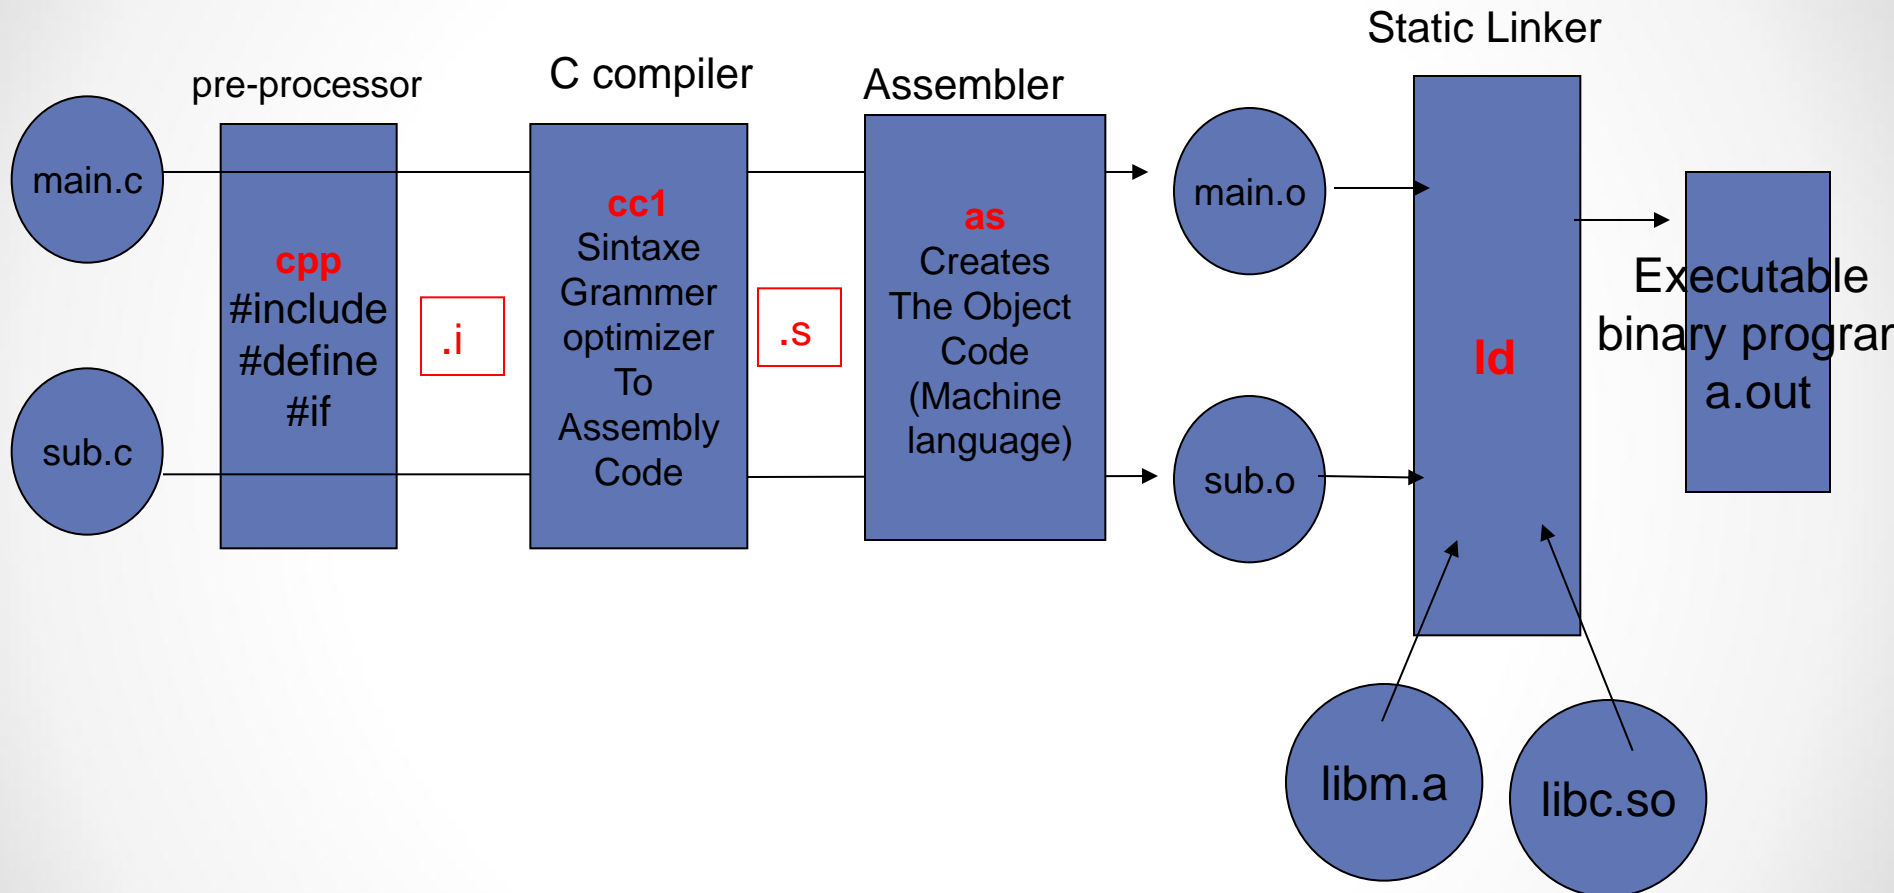
Paul Crocker

Um projecto de 2 Ficheiros

```
/*main.c*/  
#include <stdio.h>  
float sub();  
int main()  
{  
    printf("ola %f\n",sub());  
    return (0);  
}
```

```
/*sub.c*/  
#include <math.h>  
#define PI (3.14)  
const int quatro=4;  
float sub()  
{  
    #ifdef A  
        return ( sqrt (quatro*PI ) );  
    #else  
        return 1.0;  
    #endif  
}
```

Compilation Model



Static and Dynamic Libraries

Inerir código (static) ou apenas (stub)

Processo de Compilação

Tarefa	Compilação	Exercícios
Pre-Processor	<code>cc -E main.c -o main.i</code> <code>cc -E sub.c -o sub.i</code>	investigar os ficheiros <code>.i</code> less <code>main.i</code> <code>sub.i</code> O que aconteceu às linhas originais <code>#include</code> ? <code>#define</code> ? <code>const int</code> ?
Pre-Processor Options. Define a constant or Macro	<code>-D PI=3.1</code>	Defina o valor do constante PI do pre processor.
Pre-Processor Options	<code>-D A</code>	Defina o valor do constante A do pre processor. Neste caso sem valor apenas a sua existência.

Processo de Compilação

Tarefa	Compilação	Exercícios
Compiler	<code>cc -Wall -ansi -S *.i</code>	ver os ficheiros (main.s e sub.s) de assembler que são produzidos
Assembler (as)	<code>cc -c *.s</code>	i) Ver tamanhos e tipos dos ficheiros objectos. <code>ls -l *.o</code> e <code>file *.o</code> (ii) Ver symbol table <code>nm main.o</code> e <code>nm sub.o</code> (iii) Ver ficheiros usando o dissassembler <code>objdump</code> Dissambler : <code>objdump -d *.o</code> OU <code>objdump -M intel -d *.o</code>
Linker (-o calls the gnu linker ld)	<code>cc -o exemplo main.o sub.o -lm</code>	(i) Ver symbol table <code>nm exemplo</code> (2) Dissassemble <code>Objdump -d exemplo</code>

Assembler

```
main:
.LFB0:
    .cfi_startproc
    pushq    %rbp
    .cfi_def_cfa_offset 16
    .cfi_offset 6, -16
    movq     %rsp, %rbp
    .cfi_def_cfa_register 6
    movl     $0, %eax
    call     sub
    unpcklps    %xmm0, %xmm0
    cvtps2pd    %xmm0, %xmm0
    movl     $.LC0, %edi
    movl     $1, %eax
    call     printf
    movl     $0, %eax
    popq     %rbp
    .cfi_def_cfa 7, 8
    ret
```

rsp stack pointer
create the stack frame.

RETURN 0

Simplify the assembler

COMPILER EXPLORER Add... More Sponsors **intel** *PC-lint* Solid Sands Share Other Policy

C++ source #1 x x86-64 gcc 10.2 (Editor #1, Compiler #1) C++ x

A- +- v 🔍

```
1 //main.c
2 #include <stdio.h>
3 float sub ();
4 main ()
5 {
6     printf ("ola %f\n", sub ());
7     return (0);
8 }
9
```

C++

x86-64 gcc 10.2 ✓ Compiler opt

A- ⚙️ 🔍 +-

```
1 .LC0:
2     .string "ola %f\n"
3 main:
4     push    rbp
5     mov     rbp, rsp
6     call    sub()
7     pxor    xmm1, xmm1
8     cvtss2sd    xmm1, xmm0
9     movq    rax, xmm1
10    movq    xmm0, rax
11    mov     edi, OFFSET FLAT:.LC0
12    mov     eax, 1
13    call    printf
14    mov     eax, 0
15    pop     rbp
16    ret
```

Some Assembler

main:

```
pushq    %rbp          # save %rbp on the stack
movq     %rsp, %rbp    # store the value of %rsp in %rbp
movl     $0, %eax      # store the value 0 in %eax
call     sub           # set PC to sub address
popq     %rbp          # restore %rbp with the value
                        # saved on the stack
ret              # return from this function
```

Value that look like %rbp or %eax are registers → Memory that is ON the CPU.
%rbp and %rsp are special registers that refer to the base pointer and stack pointer

- PC is the program counter.

- Registers that start with “r” are 64-bits and those with “e” are 32-bits in width.

- The q suffixes on instructions refer to “quad-words” indicating that it is a 64-bit instruction.

- The l suffixes denote 32-bit instructions.

Investigate Object Files

Listing symbols from object file : nm

```
>nm main.o
000000000000000000 T main
                      U printf
                      U sub

>nm sub.o
000000000000000000 R quatro
                      U sqrt
000000000000000000 T sub
>
```

symbols

Disassembling

```
ubuntu >objdump -d main.o
```

```
main.o:   file format elf64-x86-64
```

```
Disassembly of section .text:
```

```
0000000000000000 <main>:
```

With object files

- `objdump -d`

0:	55	push %rbp
1:	48 89 e5	mov %rsp,%rbp
4:	b8 00 00 00 00	mov \$0x0,%eax
9:	e8 00 00 00 00	callq e <main+0xe>
e:	0f 14 c0	unpcklps %xmm0,%xmm0
11:	0f 5a c0	cvtps2pd %xmm0,%xmm0
14:	bf 00 00 00 00	mov \$0x0,%edi
19:	b8 01 00 00 00	mov \$0x1,%eax
1e:	e8 00 00 00 00	callq 23 <main+0x23>
23:	b8 00 00 00 00	mov \$0x0,%eax
28:	5d	pop %rbp
29:	c3	retq

Header File Directories

- Especificação de diretorias adicionais para pesquisar ficheiro de inclusão `<file.h>`
 - `-I` Additional Header File Directories (compiler)
 - Linux gcc `/usr/include` faz sempre parte da lista de diretorias pesquisados para encontrar `<file.h>`
 - `#include "/usr/include/stdio.h"` \leftrightarrow `#include </usr/include/stdio.h>`

Libraries

	Linux	Mac	Windows
Static	.a	.a	.lib
Dynamic	.so	.dylib	.dll

Notation

-lx atalho para libx.so ou libx.a (linker)
-L Additional Library Directories (linker)

Full Linker : gcc -o example main.o sub.o /usr/lib/libc.so /usr/lib/libm.so

Equivalente : gcc -o example main.o sub.o -L /usr/lib -lc -lm

... /usr/lib é quase sempre definido por defeito no linker path

Equivalente : gcc -o example main.o sub.o -lc -lm

...libc.so é incluído sempre por defeito

Equivalente : gcc -o example main.o sub.o -lm

Nota num Mac -lm não é necessário – incluído por defeito.
/usr/lib pode ser diferente p.ex /usr/lib/x86_64-linux-gnu

Static versus Dynamic

Dynamic Linking

```
>cc -o example main.o sub.o /usr/lib/x86_64-linux-gnu/libm.so
```

```
>ls -l example
```

```
-rwxrwxrwx 1 crocker crocker 8432 Mar  9 19:29 exemple
```

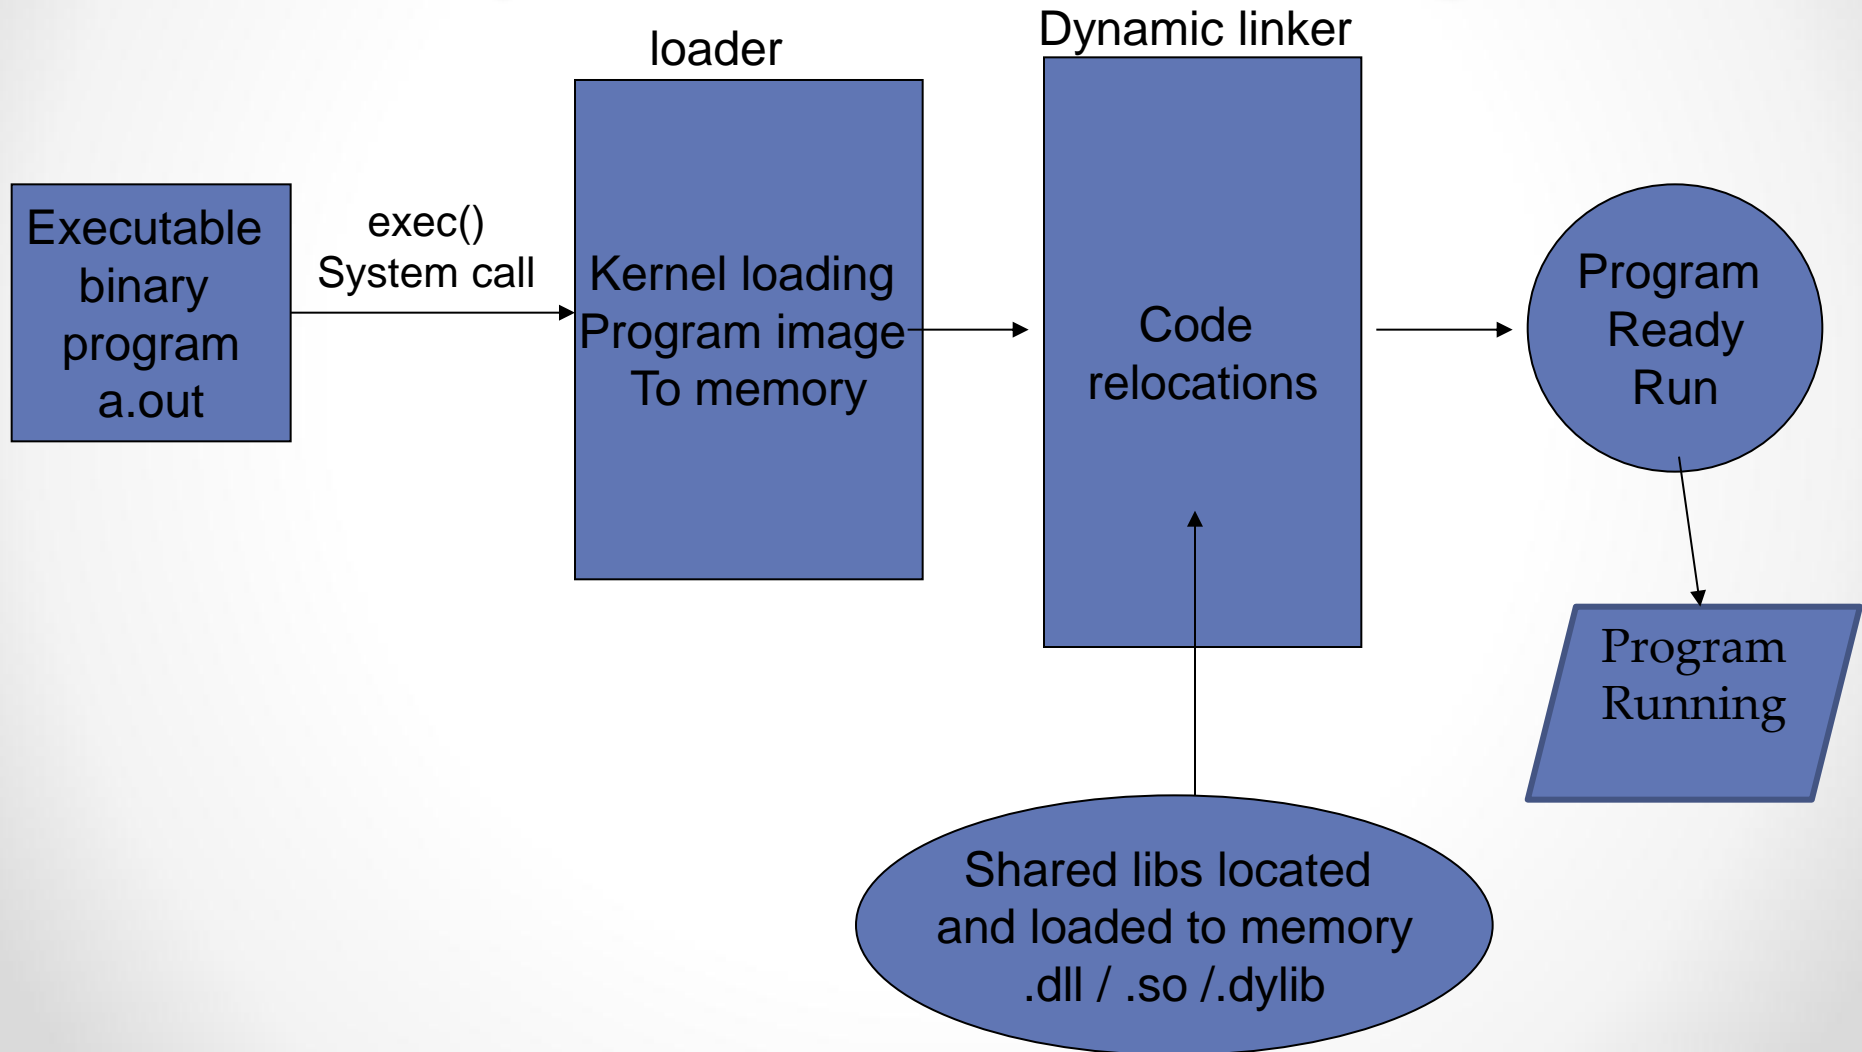
Static Linking

```
>cc -o example main.o sub.o /usr/lib/x86_64-linux-gnu/libm.a
```

```
>ls -l example
```

```
-rwxrwxrwx 1 crocker crocker 8552 Mar  9 19:29 exemple
```

Run-Time Model Dynamic Linking



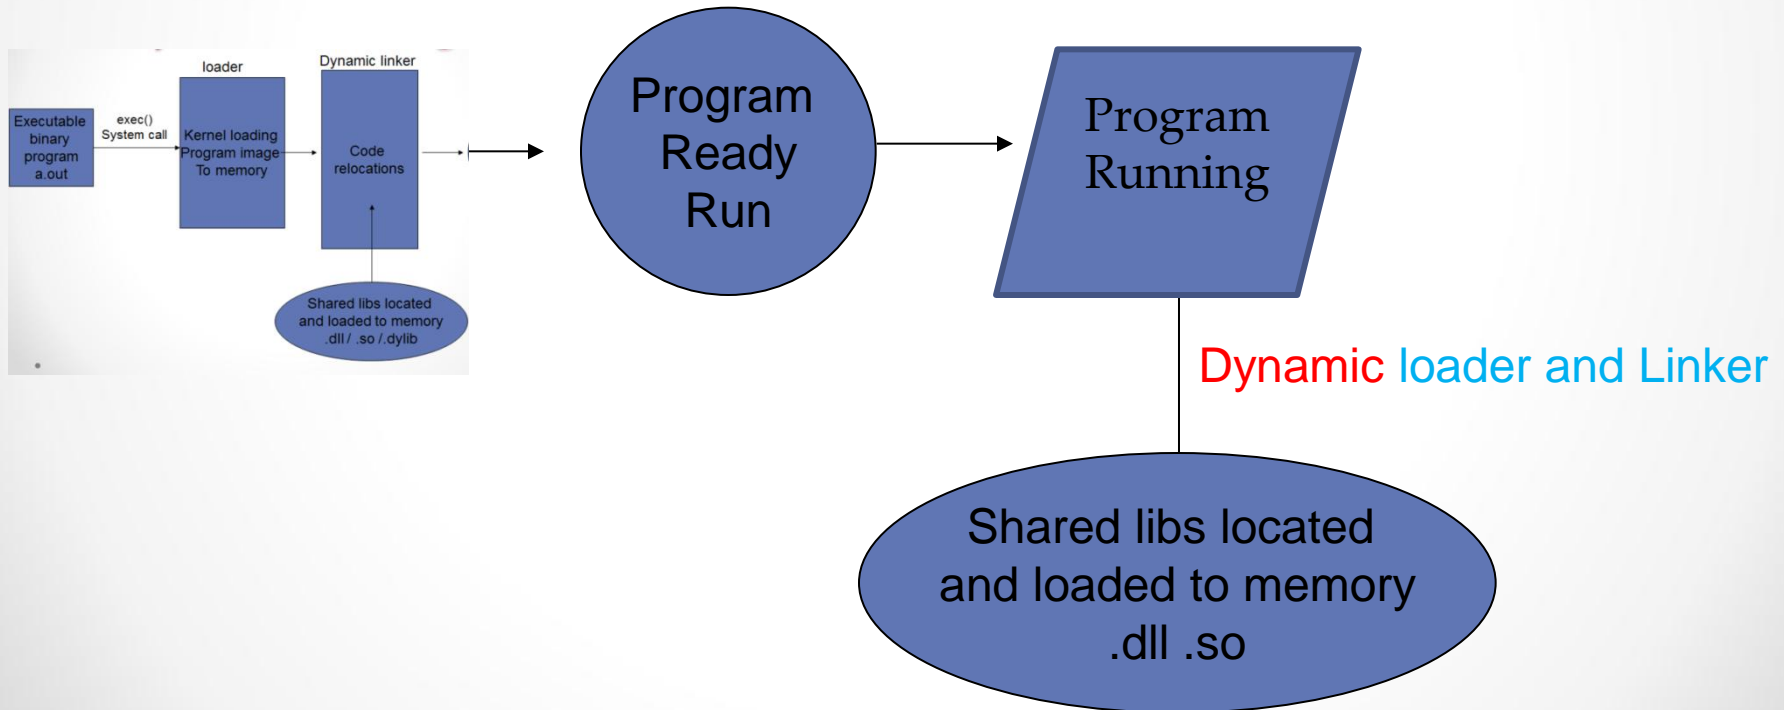
LD.SO(8) Linux Programmer's Manual

- NAME ld.so, ld-linux.so* - dynamic linker/loader
- DESCRIPTION The programs ld.so and ld-linux.so* find and load the shared libraries needed by a program, prepare the program to run, and then run it.
- Linux binaries require dynamic linking (linking at run time) unless the **-static** option was given to ld during compilation. The program ld.so handles a.out binaries, a format used long ago; ld-linux.so* handles ELF, which everybody has been using for years now.
 - Otherwise both have the same behaviour, and use the same support files and programs [ldd\(1\)](#), [ldconfig\(8\)](#) and /etc/ld.so.conf.
- The shared libraries needed by the program are searched for in various places – in particular
 - By Using the environment variable **LD_LIBRARY_PATH**.
 - etc

Run-Time Model II

Dynamic loading

Instead of Linux automatically loading and linking libraries for a given program, it's possible to **share** this control with the application itself i.e Application loads the Libraries during Run Time (NOT Linux at application startup)



Exemplo Dynamic Loading

```
#include <stdio.h> <dlfcn.h> <string.h>

void invoke_method( char *lib,
                   char *method, double argument )
{
    void *dl_handle;
    double (*func)(double);
    char *error;

    /* 1 Open the shared object */
    dl_handle = dlopen( lib, RTLD_LAZY );
    if (!dl_handle) { printf( "!!! %s\n", dlerror() );return;}

    /* 2 Resolve the symbol (method) from the object */
    func = dlsym( dl_handle, method );
    error = dlerror();
    if (error != NULL) {printf( "!!! %s\n", error );return;}

    /* 3 Call the resolved method and print the result */
    printf(" %lf\n", (*func)(argument) );

    /* 4 Close the object */
    dlclose( dl_handle );
}
```

```
#define MAX_STRING    80

int main( int argc, char *argv[] )
{
    char line[MAX_STRING+1];
    char lib[MAX_STRING+1];
    char method[MAX_STRING+1];
    double argument;

    while (1) {

        printf("> ");

        line[0]=0;
        fgets( line, MAX_STRING, stdin);

        if ( !strncmp(line, "bye", 3)) break;

        sscanf( line, "%s %s %lf", lib, method, &argument);

        invoke_method( lib, method, argument );
    }
    return 0;
}
```

Utilização

```
cc -o dynload dynload.c -ldl
cc -o dynload dynload.c -ldl
./dynload
> libm xx 2
!!! libm: cannot open shared object file: No such file or directory
> libm.so sqrt 4.0           //libm loaded to memory if necessary !
    2.000000
> libm.so cosf 0.0
    1.000000
> libm.so exp 1.0
    2.718282
> bye
```

Reference:

<http://www.ibm.com/developerworks/library/l-dynamic-libraries/>

Monitoring Run Time Execution

- Many tools and methods !!
- There are Software and Hardware monitors
- Command Line Tools
 - Linux : ps, top
 - Windows Power Shell Equivalentes
 - ps
 - `while (1) { ps | sort -desc cpu | select -first 10; sleep -seconds 2; cls }`
- Graphical Tools
- Debuggers etc.

Example: strace

- **strace** runs the specified command until it exits.
- It intercepts and records the system calls which are called by a process and the signals which are received by a process.
- The name of each system call, its arguments and its return value are printed on standard error or to the file specified with the -o option
- Each line in the trace contains the system call name, followed by its arguments and its return value.
- An example from stracing the command
 - "cat /dev/null" is:
 - `open("/dev/null", O_RDONLY) = 3`

Usando strace para ver a utilização dos system calls no printf

- printf writes to a buffer (managed by the c standard library) attached to the file : FILE *stdout.
- The actual writing to disk is done by the write() system call

```
printf("ola 1.0\n");
```

Versus

```
printf("o"); printf("l"); printf("a"); printf("\n");
```

- Quantos writes em cada caso ?

Program

```
#include <stdio.h>

int main(){
    printf("ola\n");
    printf("o"); printf("l"); printf("a"); printf("\n");
    return 0;
}
```

strace

Desligando o buffer do stdout !

- Quantas chamadas a função write() ?

```
#include <stdio.h>

int main(){
    setvbuf(stdout,NULL,_IONBF,0);
    printf("ola\n");
    printf("o"); printf("l"); printf("a"); printf("\n");
    return 0;
}
```